

<pre> ML.Uber1 import pandas as pd df = pd.read_csv('uber.csv') df.head(); <i>*pre-process the dataset</i> df.info(); df.describe(); df.isnull().sum(); print(df.isnull().values.sum()); df = df.drop(['Unnamed: 0', 'key'], axis = 1); df.head(); df = df[df.fare_amount &gt; 0]; df.shape; df.describe(); df = df[(df.passenger_count &lt;= 6) &amp; (df.passenger_count &gt; 0)]; df.head(); df = df[(df.pickup_longitude.between(-180,180,inclusive = "both")) &amp; (df.pickup_latitude.between(-90,90,inclusive = "both")) &amp; (df.dropoff_longitude.between(-180,180,inclusive = "both")) &amp; (df.dropoff_latitude.between(-90,90,inclusive = "both"))]; df.head(10); df.info(); df['pickup_datetime'] = pd.to_datetime(df['pickup_datetime']); df.info(); df.head(10); import calendar; df['year'] = df.pickup_datetime.dt.year df['month'] = df.pickup_datetime.dt.month df['weekday'] = df.pickup_datetime.dt.weekday df['hour'] = df.pickup_datetime.dt.hour; df = df.drop(['pickup_datetime'], axis=1); df.head(10); df.describe(); <i>*Remove outliers</i> df = df.reset_index(); df.head(); df = df.drop(['index'], axis = 1); df.head(10); import numpy as np; Q1 = np.percentile(df['fare_amount'], 25 ) Q3 = np.percentile(df['fare_amount'], 75 )  IQR = Q3 - Q1  print(f'IQR = {IQR}'); upper = np.where(df['fare_amount'] &gt;= (Q3 + 1.5*IQR)) lower = np.where(df['fare_amount'] &lt;= (Q1 - 1.5*IQR))  print(f'upper = {upper}') print(f'lower = {lower}'); df = df.drop(upper[0]) df = df.drop(lower[0]); df.describe(); <i>*Find Correlation</i> import seaborn as sns import matplotlib.pyplot as plt  uber_corr = df.corr() #use heatmap  plt.figure(figsize=(10,7)) sns.heatmap(uber_corr,annot=True) plt.show();  <i>*implement linear reg. &amp; random fore reg.</i> from sklearn.model_selection import train_test_split; X = df.drop('fare_amount', axis = 1) y = df['fare_amount']; </pre>	<pre> ML3GradientDescentAlgorithm import pandas as pd cur_x = 2 rate = 0.01 precision = 0.0000001 previous_step_size = 1 max_iters = 10000 iters = 0 df = lambda x: (2 * (x + 3));  while previous_step_size &gt; precision and iters &lt; max_iters:     prev_x = cur_x     cur_x = cur_x - rate * df(prev_x)     previous_step_size = abs(cur_x - prev_x)     iters = iters + 1     print(f'Iteration {iters} \n value is {cur_x} ')  print(f'The local minima occurs at {cur_x}')  ML2bankcustomer,build a neuralnetwork import pandas as pd import numpy as np; ds = pd.read_csv('Churn_Modelling.csv') ds.head(10); ds.columns; ds.shape; ds['Geography'].value_counts(normalize=True); ds = ds.drop(['RowNumber', 'CustomerId', 'Surname'], axis=1); ds.info(); ds.describe(); X = ds.iloc[:,0:10].values y = ds.iloc[:,10].values; X; from sklearn.preprocessing import LabelEncoder  print(X[:8,1], '... will now become: ')  label_X_country_encoder = LabelEncoder() X[:,1] = label_X_country_encoder.fit_transform(X[:,1]) print(X[:8,1]); print(X[:6,2], '... will now become: ')  label_X_gender_encoder = LabelEncoder() X[:,2] = label_X_gender_encoder.fit_transform(X[:,2]) print(X[:6,2]); X.shape; from sklearn.compose import ColumnTransformer from sklearn.preprocessing import OneHotEncoder  transform = ColumnTransformer([("countries", OneHotEncoder(), [1])], remainder="passthrough") X = transform.fit_transform(X); X; X = X[:,1:]; X.shape; from sklearn.model_selection import train_test_split  X_train, X_test, y_train, y_test = train_test_split(X,y, test_size = 0.2, random_state = 0) numeric_cols = ['CreditScore','Age','Tenure','Balance','NumOfProducts','EstimatedSalary']; X_train[:,np.array([2,4,5,6,7,10])]; from sklearn.preprocessing import StandardScaler </pre>
--	---

<pre> X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2); from sklearn.linear_model import LinearRegression lrmodel = LinearRegression() lrmodel.fit(X_train, y_train) lr_pred = lrmodel.predict(X_test) lr_pred;; from sklearn.ensemble import RandomForestRegressor rfmodel = RandomForestRegressor(n_estimators=100, random_state=101) rfmodel.fit(X_train, y_train) rf_pred = rfmodel.predict(X_test) rf_pred;; <i>*evaluate model</i> from sklearn.metrics import mean_squared_error;; lrmodel_rmse = np.sqrt(mean_squared_error(y_test, lr_pred)) rfmodel_rmse = np.sqrt(mean_squared_error(y_test, rf_pred))  print(f'Linear Regression RMSE = {lrmodel_rmse}') print(f'Random Forest RMSE = {rfmodel_rmse}'); from sklearn.metrics import r2_score lrmodel_r2 = r2_score(y_test, lr_pred) rfmodel_r2 = r2_score(y_test, rf_pred)  print(f'Linear Regression R2 = {lrmodel_r2}') print(f'Random Forest R2 = {rfmodel_r2}');  MLSKMeansClustering import pandas as pd import numpy as np import seaborn as sns import matplotlib.pyplot as plt;; from sklearn.cluster import KMeans, k_means from sklearn.decomposition import PCA;; df = pd.read_csv('/content/sales_data_sample.csv');; df.head();; df.shape;; df.describe();; df.info();; df.isnull().sum();; df.dtypes;; df_drop = ['ADDRESSLINE1', 'ADDRESSLINE2', 'STATUS', 'POSTALCODE', 'CITY', 'TERRITORY', 'PHONE', 'STATE', 'CONTACTFIRSTNAME', 'CONTACTLASTNAME', 'CUSTOMERNAME', 'ORDERNUMBER'] df = df.drop(df_drop, axis=1);; df.isnull().sum();; df.dtypes;; <i>*checking the categorical columns</i> df['COUNTRY'].unique();; df['PRODUCTLINE'].unique();; df['DEALSIZE'].unique();; productline = pd.get_dummies(df['PRODUCTLINE']) Dealsize = pd.get_dummies(df['DEALSIZE']);; df = pd.concat([df, productline, Dealsize], axis = 1);; df_drop = ['COUNTRY', 'PRODUCTLINE', 'DEALSIZE'] #Dropping Country too as there are alot of countries. df = df.drop(df_drop, axis=1);; df['PRODUCTCODE'] = pd.Categorical(df['PRODUCTCODE']).codes;; df.drop('ORDERDATE', axis=1, inplace=True);; df.dtypes;; distortions = [] K = range(1,10) for k in K:     kmeanModel = KMeans(n_clusters=k)     kmeanModel.fit(df)     distortions.append(kmeanModel.inertia_);; </pre>	<pre> sc=StandardScaler() X_train[:,np.array([2,4,5,6,7,10])] = sc.fit_transform(X_train[:,np.array([2,4,5,6,7,10])]) X_test[:,np.array([2,4,5,6,7,10])] = sc.transform(X_test[:,np.array([2,4,5,6,7,10])]);; X_train[0];; from sklearn.preprocessing import StandardScaler  sc=StandardScaler() X_train = sc.fit_transform(X_train) X_test = sc.transform(X_test);; X_train;; X_train.shape;; from tensorflow.keras.models import Sequential  # Initializing the ANN classifier = Sequential();; from tensorflow.keras.layers import Dense classifier.add(Dense(activation = 'relu', input_dim = 11, units=16, kernel_initializer='uniform'));; classifier.add(Dense(8, activation='relu', kernel_initializer='uniform'));; classifier.add(Dense(1, activation = 'sigmoid', kernel_initializer='uniform'));; 192-11*16;; classifier.summary();; classifier.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy']);; classifier.summary();; classifier.fit(X_train, y_train, validation_data=(X_test,y_test), epochs=20);; y_pred = classifier.predict(X_test) print(y_pred);; y_pred = (y_pred &gt; 0.5) print(y_pred);; from sklearn.metrics import confusion_matrix,classification_report  cm1 = confusion_matrix(y_test, y_pred) print(cm1);; print(classification_report(y_test, y_pred));; accuracy_model1 = ((cm1[0][0]+cm1[1][1])*100)/(cm1[0][0]+cm1[1][1]+cm1[0][1]+ cm1[1][0]) print (accuracy_model1, '% of testing data was classified correctly');; classifier.summary();; classifier.compile(optimizer='adam', loss = 'binary_crossentropy', metrics=['accuracy']);; classifier.fit(X_train, y_train, validation_data=(X_test,y_test), epochs=20, batch_size=32);; y_pred = classifier.predict(X_test) print(y_pred);; y_pred = (y_pred &gt; 0.5) print(y_pred);; cm2 = confusion_matrix(y_test, y_pred) print(cm2);; cm2 = classification_report(y_test, y_pred) print(cm2);; </pre>
--	---

<pre> plt.figure(figsize=(16,8)) plt.plot(K, distortions, 'bx-') plt.xlabel('k') plt.ylabel('Distortion') plt.title('The Elbow Method showing the optimal k') plt.show(); <i>*numb. Of k increases inertia decreases</i> X_train = df.values;; X_train.shape;; model = KMeans(n_clusters=3,random_state=2) model = model.fit(X_train) predictions = model.predict(X_train);; unique,counts = np.unique(predictions,return_counts=True);; counts = counts.reshape(1,3);; counts_df = pd.DataFrame(counts,columns=['Cluster1','Cluster2','Cluster3']);; counts_df.head();; pca = PCA(n_components=2);; reduced_X = pd.DataFrame(pca.fit_transform(X_train),columns=['PCA1','PCA2']);; reduced_X.head();; plt.figure(figsize=(14,10)) plt.scatter(reduced_X['PCA1'],reduced_X['PCA2']);; model.cluster_centers_;; reduced_centers = pca.transform(model.cluster_centers_);; reduced_centers;; plt.figure(figsize=(14,10)) plt.scatter(reduced_X['PCA1'],reduced_X['PCA2']) plt.scatter(reduced_centers[:,0],reduced_centers[:,1],color='black',mark er='x',s=300);; reduced_X['Clusters'] = predictions;; reduced_X.head();; plt.figure(figsize=(14,10))  plt.scatter(reduced_X[reduced_X['Clusters'] == 0].loc[:, 'PCA1'],reduced_X[reduced_X['Clusters'] == 0].loc[:, 'PCA2'],color='slateblue') . plt.scatter(reduced_X[reduced_X['Clusters'] == 1].loc[:, 'PCA1'],reduced_X[reduced_X['Clusters'] == 1].loc[:, 'PCA2'],color='springgreen') . plt.scatter(reduced_X[reduced_X['Clusters'] == 2].loc[:, 'PCA1'],reduced_X[reduced_X['Clusters'] == 2].loc[:, 'PCA2'],color='indigo') .  plt.scatter(reduced_centers[:,0],reduced_centers[:,1],color='black',mark er='x',s=300);; </pre>	<pre> ML4KNN import pandas as pd import numpy as np import seaborn as sns import matplotlib.pyplot as plt %matplotlib inline import warnings warnings.filterwarnings('ignore') from sklearn.model_selection import train_test_split from sklearn.svm import SVC from sklearn import metrics;; df=pd.read_csv('/content/diabetes.csv');; df.columns;; df.isnull().sum();; X = df.drop('Outcome',axis = 1) y = df['Outcome'];; from sklearn.preprocessing import scale X = scale(X) # split into train and test X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 42);; from sklearn.neighbors import KNeighborsClassifier knn = KNeighborsClassifier(n_neighbors=7)  knn.fit(X_train, y_train) y_pred = knn.predict(X_test);; print("Confusion matrix: ") cs = metrics.confusion_matrix(y_test,y_pred) print(cs);; print("Accuracy ",metrics.accuracy_score(y_test,y_pred));; total_misclassified = cs[0,1] + cs[1,0] print(total_misclassified) total_examples = cs[0,0]+cs[0,1]+cs[1,0]+cs[1,1] print(total_examples) print("Error rate",total_misclassified/total_examples) print("Error rate ",1-metrics.accuracy_score(y_test,y_pred));; print("Precision score",metrics.precision_score(y_test,y_pred));; print("Recall score ",metrics.recall_score(y_test,y_pred));; print("Classification report ",metrics.classification_report(y_test,y_pred));; </pre>
---	---

